

## Comparison Sorts III

continued

---

CS 311 Data Structures and Algorithms

Lecture Slides

Wednesday, March 6, 2013

Chris Hartman

Department of Computer Science

University of Alaska Fairbanks

cmhartman@alaska.edu

Based on material by Glenn G. Chappell

© 2005–2009 Glenn G. Chappell

# Unit Overview

## Algorithmic Efficiency & Sorting

---

### Major Topics

- ✓ ■ Introduction to Analysis of Algorithms
- ✓ ■ Introduction to Sorting
- ✓ ■ Comparison Sorts I
- ✓ ■ More on Big- $O$
- ✓ ■ The Limits of Sorting
- ✓ ■ Divide-and-Conquer
- ✓ ■ Comparison Sorts II
- (part) ■ Comparison Sorts III
  - Radix Sort
  - Sorting in the C++ STL

# Review

## Introduction to Analysis of Algorithms

---

### Efficiency

- General: using few resources (time, space, bandwidth, etc.).
- Specific: fast (time).

### Analyzing Efficiency

- Determine how the **size of the input** affects running time, measured in **steps**, in the **worst case**.

**Scalable:** works well with large problems.

	Using Big-O	In Words	
	$O(1)$	Constant time	
	$O(\log n)$	Logarithmic time	
	$O(n)$	Linear time	
	$O(n \log n)$	Log-linear time	
..... Probably not scalable ↓	$O(n^2)$	Quadratic time	
	$O(b^n)$ , for some $b > 1$	Exponential time	

Cannot read all of input ↑

Faster ↑

Slower ↓

## Review


### Introduction to Sorting — Basics, Analyzing

**Sort:** Place a collection of data in order.

**Key:** The part of the data item used to sort.

**Comparison sort:** A sorting algorithm that gets its information by comparing items in pairs.

A **general-purpose comparison sort** places no restrictions on the size of the list or the values in it.



Five criteria for analyzing a general-purpose comparison sort:

- (Time) Efficiency
  - Requirements on Data
  - Space Efficiency
  - Stability
  - Performance on Nearly Sorted Data
- In-place** = no large additional space required.
- Stable** = never changes the order of equivalent items.
1. All items close to proper places, OR  
2. few items out of order.

## Review

### Introduction to Sorting — Overview of Algorithms

---

There is no *known* sorting algorithm that has all the properties we would like one to have.

We will examine a number of sorting algorithms. Most of these fall into two categories:  $O(n^2)$  and  $O(n \log n)$ .


- Quadratic-Time [ $O(n^2)$ ] Algorithms
  - ✓ ■ Bubble Sort
  - ✓ ■ Insertion Sort
  - (part) ■ Quicksort
    - Treesort (later in semester)
- Log-Linear-Time [ $O(n \log n)$ ] Algorithms
  - ✓ ■ Merge Sort
    - Heap Sort (mostly later in semester)
    - Introsort (not in the text)
- Special Purpose — Not Comparison Sorts
  - Pigeonhole Sort
  - Radix Sort

## Review

### Comparison Sorts I — Insertion Sort

---

**Insertion Sort** repeatedly does this:



#### Analysis

- Efficiency:  $O(n^2)$ . Average case same. ☹️
- Requirements on data: Works for Linked Lists, etc. 😊
- Space Efficiency: In-place. 😊
- Stable: Yes. 😊
- Performance on Nearly Sorted Data:  $O(n)$  for both kinds. 😊

#### Notes

- Too slow for general-purpose use.
- Works well on **nearly sorted data** and **small lists**.
- Thus, used as **part of other algorithms**.

## Review

### More on Big- $O$ , The Limits of Sorting

---

Three ways to talk about how fast a function grows.

$g(n)$  is:

- $O(f(n))$  if  $g(n) \leq k \times f(n) \dots$
- $\Omega(f(n))$  if  $g(n) \geq k \times f(n) \dots$
- $\Theta(f(n))$  if both of the above are true.
  - Possibly with different values of  $k$ .

In an algorithmic context,  $g(n)$  might be the max number of steps required by some algorithm when given input of size  $n$ , or the max amount of additional space required.

---

**Fact.** Every general-purpose comparison sort does  $\Theta(n \log n)$  comparisons in the worst case.

## Review

### Divide-and-Conquer

---

A common algorithmic strategy is called **divide-and-conquer**: split the input into pieces, and handle these with recursive calls.

If an algorithm using divide-and-conquer splits the input into **nearly equal-sized parts**, then we can analyze it using the **Master Theorem**.

- $b$  is the number of nearly equal-sized parts. Important!
- $b^k$  is the number of recursive calls. Find  $k$ .
- $f(n)$  is the amount of extra work done (number of steps).
  - Hopefully,  $f(n)$  looks like  $n$  raised to some power.
- If the power is less than  $k$ , then the algorithm is  $O(n^k)$ .
- If the power is  $k$ , then the algorithm is  $\Theta(n^k \log n)$ .



## Review

### Comparison Sorts II — Merge Sort

**Merge Sort** splits the data in half, recursively sorts each half, and then merges the two.

#### Stable Merge


- Linear time, stable.
- In-place for Linked List. Uses buffer [ $O(n)$  space] for array.

#### Analysis

- Efficiency:  $O(n \log n)$ . Average same. 😊
- Requirements on data: Works for Linked Lists, etc. 😊
- Space Efficiency:  $O(\log n)$  space for Linked List. Can eliminate recursion to make this in-place.  $O(n)$  space for array. 😊/😊/😞
- Stable: Yes. 😊
- Performance on Nearly Sorted Data: Not better or worse. 😊

#### Notes

- Practical & often used.
- Fastest known for (1) stable sort, (2) sorting a Linked List.
- Good standard for judging sorting algorithms



## Review

### Comparison Sorts III — Quicksort: Introduction, Partition


**Quicksort** is another divide-and-conquer algorithm. Procedure:

- Choose a list item (the **pivot**).
- Do a **Partition**: put items less than the pivot before it, and items greater than the pivot after it.
- Recursively sort two sublists: items before pivot, items after pivot.

We did a simple pivot choice: the first item. Later, we improve this.

Fast Partition algorithms are in-place, but not stable.

- Note: In-place Partition does not give us an in-place Quicksort. Quicksort uses memory for recursion.



## Review

### Comparison Sorts III — Better Quicksort: Problem

---

Quicksort has a problem.

- In the worst case, the pivot is chosen poorly.
- Thus: linear recursion depth, and so Quicksort is  $O(n^2)$ . ☹
- And the worst case happens when the list is **already sorted**!

However, Quicksort's **average-case** time is very fast.

- This is  $O(n \log n)$  and typically significantly faster than Merge Sort.

Quicksort is *usually* very fast; thus, people want to use it.

- So we try to figure out how to make it better.
- We look at three of the best optimizations ...

## Review


### Comparison Sorts III — Better Quicksort: Opt. 1 = Pivot Selection

Choose the pivot using **median-of-3**.

- Look at 3 items in the list: first, middle, last.
- Let the pivot be the one that is between the other two (by <).


#### Basic Quicksort

Initial State:



Pivot


After Partition:



↖ ↗  
Recursively Sort


#### Quicksort with Median-of-3 Pivot Selection

Initial State:



Pivot

After Partition:



↖ ↗  
Recursively Sort

This gives acceptable performance on most nearly sorted data.

- But it is still  $O(n^2)$ .

## Comparison Sorts III

### Better Quicksort — Optimization 2: Tail-Recursion Elimination

---

How much additional space does Quicksort use?

- Partition is in-place and Quicksort uses few local variables.
- However, Quicksort is recursive.
- Quicksort's additional space usage is thus proportional to its recursion depth ...
- ... which is linear. Additional space:  $O(n)$ . ☹

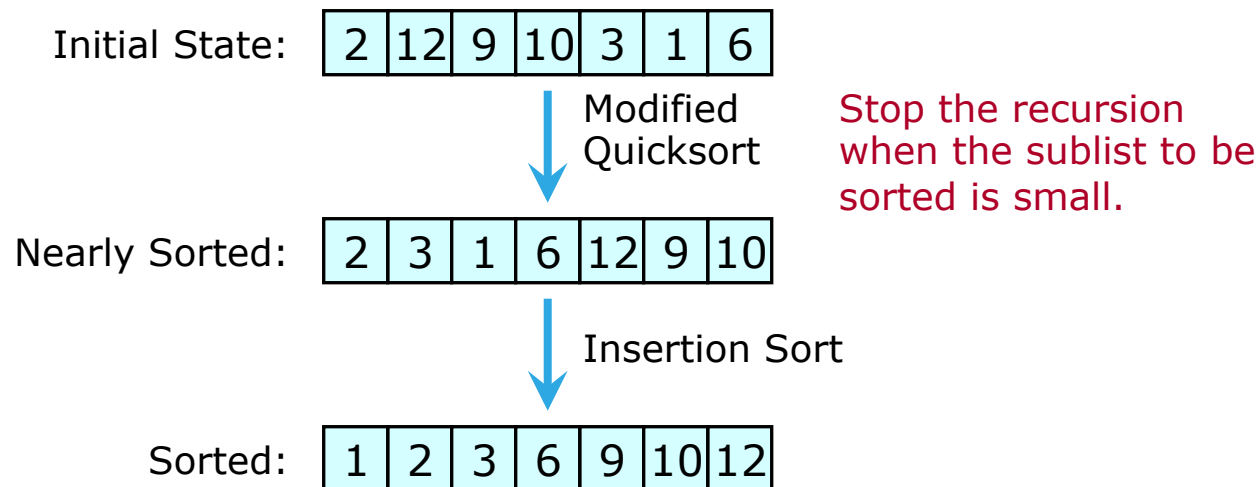
We can significantly improve this:

- Do the **larger** of the two recursive calls last.
- Do tail-recursion elimination on this final recursive call.
- Result: Recursion depth & additional space usage:  $O(\log n)$ . ☺
  - And this additional space need not hold any data items.

### Better Quicksort — Optimization 3: Finishing with Insertion Sort

#### Another Speed-Up: Finish with Insertion Sort

- Stop Quicksort from going to the bottom of its recursion. We end up with a nearly sorted list.
- Finish sorting this list using one call to Insertion Sort.
- This is generally faster\*, but still  $O(n^2)$ .
- Note: This is not the same as using Insertion Sort for small lists.



\*I have read that this tends to adversely affect the number of cache hits.

## Comparison Sorts III

### Better Quicksort — Rewrite It

---

#### TO DO

- Rewrite our Quicksort to do:
  - Median-of-3 pivot selection.
  - Tail-recursion elimination for reduced recursion depth.
  - Finishing with Insertion Sort.

## Comparison Sorts III

### Better Quicksort — Needed?

---

We want an algorithm that:

- Is as fast as Quicksort on the average.
- Has reasonable [ $O(n \log n)$ ] worst-case performance.

But for over three decades no one found one.

Some said (and some still say) “Quicksort’s bad behavior is very rare; ignore it.”

- I suggest to you that this is not a good way to think.
- Sometimes bad worst-case behavior is okay; sometimes it is not.
  - Know what is important in the situation you are addressing.
  - Also, understand that your software can end up being used in other situations.
  - Lastly, remember that on the Web, there are **malicious users**.
- From a former version of the Wikipedia article on Quicksort (retrieved 18 Oct 2006; the statements below were removed on 19 Jan 2007):

The worst-case behavior of quicksort is not merely a theoretical problem. When quicksort is used in web services, for example, it is possible for an attacker to deliberately exploit the worst case performance and choose data which will cause a slow running time or maximize the chance of running out of stack space.

However, in 1997, a solution was finally published. We discuss this shortly. But first, we analyze Quicksort.



## Comparison Sorts III

### Better Quicksort — Analysis of Quicksort

---

#### Efficiency ☹

- Quicksort is  $O(n^2)$ .
- Quicksort has a **very** good  $O(n \log n)$  average-case time. 😊😊

#### Requirements on Data ☹

- Non-trivial pivot-selection algorithms (median-of-3 and others) are only efficient for random-access data.

#### Space Usage ☹

- Quicksort uses space for recursion.
  - Additional space:  $O(\log n)$ , if clever tail-recursion elimination is done.
  - Even if **all** recursion is eliminated,  $O(\log n)$  additional space is still used.
  - This additional space need not hold any data items.

#### Stability ☹

- Efficient versions of Quicksort are not stable.

Unlike  
Merge Sort

#### Performance on Nearly Sorted Data ☹

- An unoptimized Quicksort is **slow** on nearly sorted data:  $O(n^2)$ .
- Quicksort + median-of-3 is  $O(n \log n)$  on most nearly sorted data.

## Comparison Sorts III

### Introsort — “Introspection”

---

In 1997, David Musser introduced a simple algorithmic-design idea.

- For a number of problems, there are known algorithms with very good average-case performance and very bad worst-case performance.
- Quicksort is the best known of these, but there are others.
- Musser suggests keeping track of an algorithm’s performance. If it is not doing well, switch to a different algorithm that has reasonably good worst-case performance.
- Musser calls this technique **introspection**, since an algorithm is examining itself.

The most important application is to sorting.

- Now we can eliminate the bad behavior of Quicksort.

## Comparison Sorts III

### Introsort — Heap Sort Preview

---

This is a preview of a sorting algorithm to be covered later.

Later in the semester, we will study the “Priority Queues”, generally implemented via a data structure known as a “Heap”.

- In a normal Queue, we insert items, and then remove them in the same order (first-in-first-out).
- In a **Priority Queue**, each item has a “priority”. Items come out in order of their priority.

Set an item’s priority equal to its numerical value, and items come out in sorted order.

- So: Make a Heap and then remove all items from it, in numerical order.
- This sorting algorithm is called **Heap Sort**.

Important Facts about Heap Sort

- Log-linear time.
- In-place.
- Requires random-access data.
- Can be modified to handle problems that are more general than simple comparison sorting. For example, we can allow new items to be added during the sorting process.
- Used as part of a very fast sorting algorithm called “Introsort”. Read on ...

## Comparison Sorts III

### Introsort — Description

---

Recall: Quicksort does a linear time operation (Partition), then calls itself recursively.

- If the recursion depth is around  $\log n$ , then it uses  $O(n \log n)$  steps.
  - Count both sub-lists as recursive calls. Ignore the tail-recursion trick.
- Thus, Quicksort is slow only **when the recursion gets too deep**.

Apply introspection:

- Do optimized Quicksort, but keep track of the recursion depth.
- If the depth exceeds some threshold ( $k \log n$ , for some  $k$ ), switch to Heap Sort for the current sublist being sorted.
  - Musser suggested a threshold of  $2 \log_2 n$ .

The resulting algorithm is called **Introsort**.

Musser's 1997 paper discusses the speed-ups we have covered:

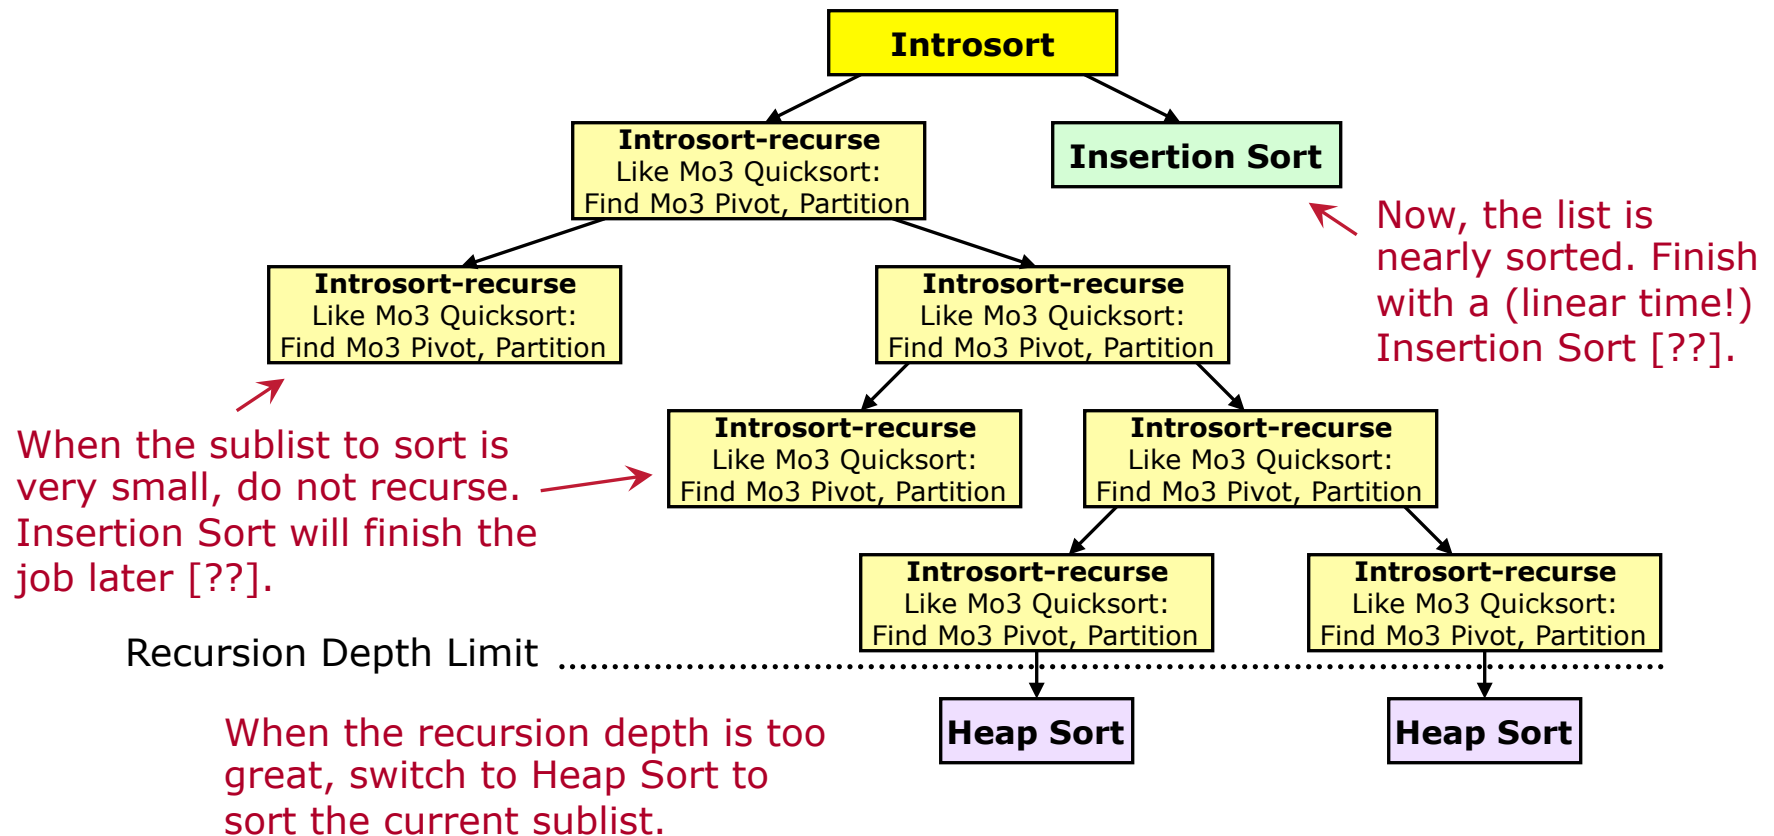
- Use the median-of-3 rule for pivot selection.
- Stop the recursion prematurely, and finish with Insertion Sort.
  - Maybe. This can adversely affect cache performance.
- However, it is no longer necessary to handle the larger and smaller recursive calls differently, since the recursion-depth limit already makes sure that excessive recursive calls are not made.

## Comparison Sorts III

### Introsort — Diagram

Here is an illustration of how Introsort works.

- In practice, the recursion will be deeper than this.
- The Insertion-Sort call *might* not be done, due to its effect on cache hits.



## Comparison Sorts III

### Introsort — Analysis

---

#### Efficiency 😊😊

- Introsort is  $O(n \log n)$ .
- Introsort also has an average-case time of  $O(n \log n)$  [of course].
  - Its average-case time is just as good as Quicksort. 😊😊

#### Requirements on Data 😞

- Introsort requires random-access data.

#### Space Usage 😐

- Introsort uses space for recursion (or simulated recursion).
  - Additional space:  $O(\log n)$  — even if all recursion is eliminated.
  - This additional space need not hold any data items.

#### Stability 😞

- Introsort is not stable.

#### Performance on Nearly Sorted Data 😐

- Introsort is not significantly faster or slower on nearly sorted data.

## Comparison Sorts III

### Introsort — Comments

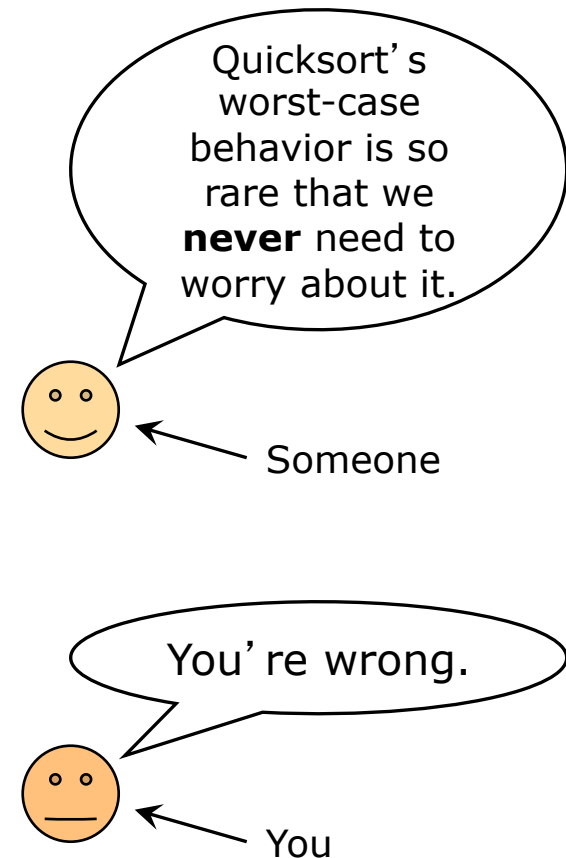
---

In general, when you want speed (i.e., all the time), Introsort is usually the algorithm to use.

#### Some Exceptions

- When you need a stable sort.
- When sorting non-random-access data.
- When you expect the data to be nearly sorted.
- When memory is limited.
  - On some embedded systems, maybe?
  - Sorting a list that does not fit into memory.
- When data are accessed over a slow connection.
  - Sorting data accessed over a network?
- When the problem you want to solve is not exactly comparison sorting.
- When your sort may be used in multiple applications.

More about some of these later in the semester.  
If someone tells you that Quicksort's worst-case behavior is so rare that we **never** need to worry about it, tell them they're wrong.



## Comparison Sorts III

### When is it Best?

---

Algorithm	When This Algorithm is the <i>Best One</i>
Bubble Sort	Never
Insertion Sort	<ul style="list-style-type: none"><li>■ For small lists</li><li>■ When you are guaranteed nearly sorted data</li></ul>
Merge Sort	<ul style="list-style-type: none"><li>■ When stability is needed</li><li>■ For special data types, especially Linked Lists</li></ul>
Heap Sort	<p>In certain special situations:</p> <ul style="list-style-type: none"><li>■ When a list is operated on during the sorting process</li><li>■ When you only care about the ordering of part of a list</li><li>■ Etc. (more about this later in the semester)</li></ul>
Quicksort	Never
Introsort	Most of the time (if you do not care about stability, data accessed via slow connections, sequential-access data, ...)

Now, what if (say) Quicksort is written for you, but nothing else is? Should you write your own? Maybe. It depends on the situation. **Think!**