

Node-Based Structures

More on Linked Lists

CS 311 Data Structures and Algorithms
Lecture Slides
Friday, March 29, 2013
Chris Hartman

Department of Computer Science
University of Alaska Fairbanks
cmhartman@alaska.edu
Based on material by Glenn G. Chappell
© 2005–2009 Glenn G. Chappell

Unit Overview

Handling Data & Sequences

Major Topics

- ✓ ■ Data abstraction
- ✓ ■ Introduction to Sequences
- ✓ ■ Smart arrays
 - ✓ ■ Array interface
 - ✓ ■ Basic array implementation
 - ✓ ■ Exception safety
 - ✓ ■ Allocation & efficiency
 - ✓ ■ Generic containers
- Linked Lists
 - Node-based structures
 - More on Linked Lists
- Sequences in the C++ STL
- Stacks
- Queues

Review

Exception Safety [1/2]

Safety: Does function ever signal an error condition, and if so:

- Are data left in a usable state?
- Do we know something about that state?
- Are resource leaks avoided?

Basic Guarantee



Minimum
standard

- Data remain in a usable state, and resources are never leaked, even in the presence of exceptions.

Each guarantee
includes the
previous one.

Strong Guarantee



Preferred

- If the operation throws an exception, then it makes no changes that are visible to the client code.

No-Throw Guarantee



Required in some
special situations

- The operation never throws an exception.

Review

Exception Safety [2/2]

To make sure code is exception-safe:

- Look at **every** place an exception might be thrown.
- For each, make sure that, if an exception is thrown, either
 - we terminate normally and meet our postconditions, or
 - we throw and meet our guarantees.

A bad design can force us to be unsafe.

- Thus, good design is part of exception safety.
- An often helpful idea is that every module has exactly one well defined responsibility (the **Single Responsibility Principle**).
- In particular: A non-const member function should not return an object by value.

Review

Allocation & Efficiency [1/2]

An operation is **amortized constant time** if k operations require $O(k)$ time.

- Thus, over *many consecutive operations*, the operation averages constant time.
- Not the same as constant-time average case (which averages over *all possible inputs*)
- Quintessential amortized-constant-time operation: insert-at-end for a well written smart array.
- Amortized constant time is not something we can easily compare with (say) logarithmic time.

Review

Allocation & Efficiency [2/2]

Improving **SmArray**

- Our original design did not allow for efficient insert-at-end.
 - Reallocate-and-copy would happen every time.
- The revised design had three data members: size, capacity, and the array pointer.
- Having a “capacity” member allows us to keep a record of how much memory is allocated. Then we can allocate extra so that we do not need to reallocate every time.
- Result: We can do amortized-constant-time insert-at-end.

Review

Generic Containers [1/2]

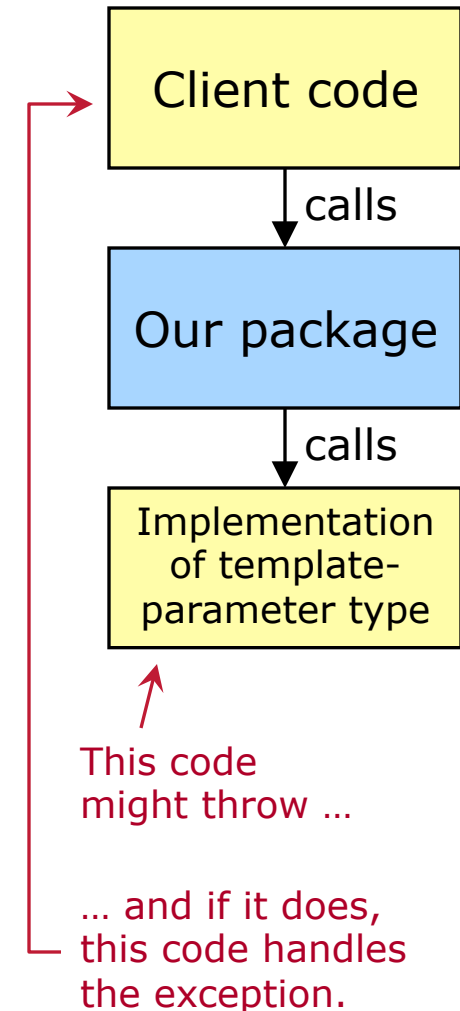
A **generic container** is a container that can hold a client-specified data type.

- In C++ we usually implement a generic container using a **class template**.

A function that allows exceptions thrown by a client's code to propagate unchanged, is said to be **exception-neutral**.

When exception-neutral code calls a client-provided function that may throw, it does one of two things:

- Call the function outside a try block, so that any exceptions terminate our code immediately.
- Or, call the function inside a try block, then catch all exceptions, do any necessary clean-up, and re-throw.



Review

Generic Containers [2/2]

We can use catch-all, clean-up, re-throw to get both exception safety and exception neutrality.

```
arr = new MyType[size];  
try  
{  
    std::copy(a, a+size, arr);  
}  
catch (...)  
{  
    delete [] arr;  
    throw;  
}
```

← Called outside any `try` block. If this fails, we exit immediately, throwing an exception.

← Called inside a `try` block. If this fails, we need to deallocate the array before exiting.

← This helps us meet the Basic Guarantee (also the Strong Guarantee if this function does nothing else).

← This makes our code exception-neutral.

Node-Based Structures

Introduction [1/2]


Our fundamental building block for data structures has been the **array**.



- Items are stored in contiguous memory locations.
- Look-up operations are usually very fast.
- Operations that require rearrangement (insert, remove, etc.) can be slow.

Now we begin looking at data structures built out of **nodes**.

- A *node* is generally a small block of memory that is referenced via a pointer, and which may reference other nodes via pointers.




- Node-based structures do not necessarily store data in contiguous memory.
 - Memory-management changes significantly.
- To find a node, we follow a chain of pointers. Look-up can be slow.
- Operations that require rearrangement can be very fast.

Node-Based Structures


Introduction [2/2]

When we draw pictures of node-based data structures, the positions of nodes in the picture usually have nothing to do with their positions in memory.

For example, if a structure is stored like this ...




... then we might draw it like this:





Node-Based Structures Implementation — Storage Options

We normally store nodes in separately allocated blocks of memory.



However, we might allocate an **array of nodes**.

- Replace pointers with array indices, if desired.



Both of these
have the
above **logical**
structure.

This is still a node-based data structure.

- It still has most of the pros & cons of node-based structures.
- The main differences involve **memory management**: who does it & when it gets done.

Node-Based Structures Implementation — Classes

When we define a class to implement a node-based data structure, we may wish to write several classes:

- The main **container class**, representing the structure as a whole.
- A class representing a **node**.
 - This might be a private member of the main class.
 - Typically client code does not deal with this class.
 - If there are different kinds of nodes, then we might want several node classes. Sometimes inheritance can be helpful here.
- Possibly an **iterator** class.
 - Iterators to node-based structures are almost never pointers, because `operator++`, for example, needs to go to the next **logical** node, not the next **physical** memory location.

Despite the multiple classes being defined, we are implementing only a single package.

- Thus, multiple header & source files are generally not necessary.
- We can make all of these classes friends without breaking encapsulation.

Node-Based Structures

Implementation — Pointers & Ownership

Think of nodes as resources to be owned & managed.

- Who owns them?
 - Always document ownership (here: in the class invariants).
- Internal pointers in a node-based structure will usually be owning pointers.
 - A node is typically owned by the node that points to it.
 - Thus, a node's destructor should free all nodes that it points to.

This can make destroying a node-based structure **easy**.

- Each node is responsible for destroying the nodes it owns.
- Thus, to destroy the whole structure, all we need to do is destroy the nodes that are not owned by other nodes.
- And there is usually just one of these.

More on Linked Lists

Refresher [1/2]


Earlier in the semester, we looked briefly at the simplest node-based structure: a **Linked List**.

- Like an array, a Linked List is a structure for storing a sequence of items.

Array

3	1	5	3	5	2
---	---	---	---	---	---

- A Linked List is composed of **nodes**. Each has a single data item and a pointer to the next node.



- These pointers are the **only** way to find the next data item. Thus, unlike an array, we cannot quickly skip to (say) the 100th item in a Linked List. Nor can we quickly find the previous item.
- A Linked List is a one-way sequential-access data structure. Thus, its natural iterator is a **forward iterator**, which has only the ++ operator.

More on Linked Lists


Refresher [2/2]

Why not always use (smart) arrays?

- One important reason: we can often insert and remove much faster with a Linked List.


Inserting

- Inserting an item at a given position in an array is slow-ish.
- Inserting an item at a given position (think “iterator”) in a Linked List is very fast.
- Example: insert a “7” after the bold node.



Removing

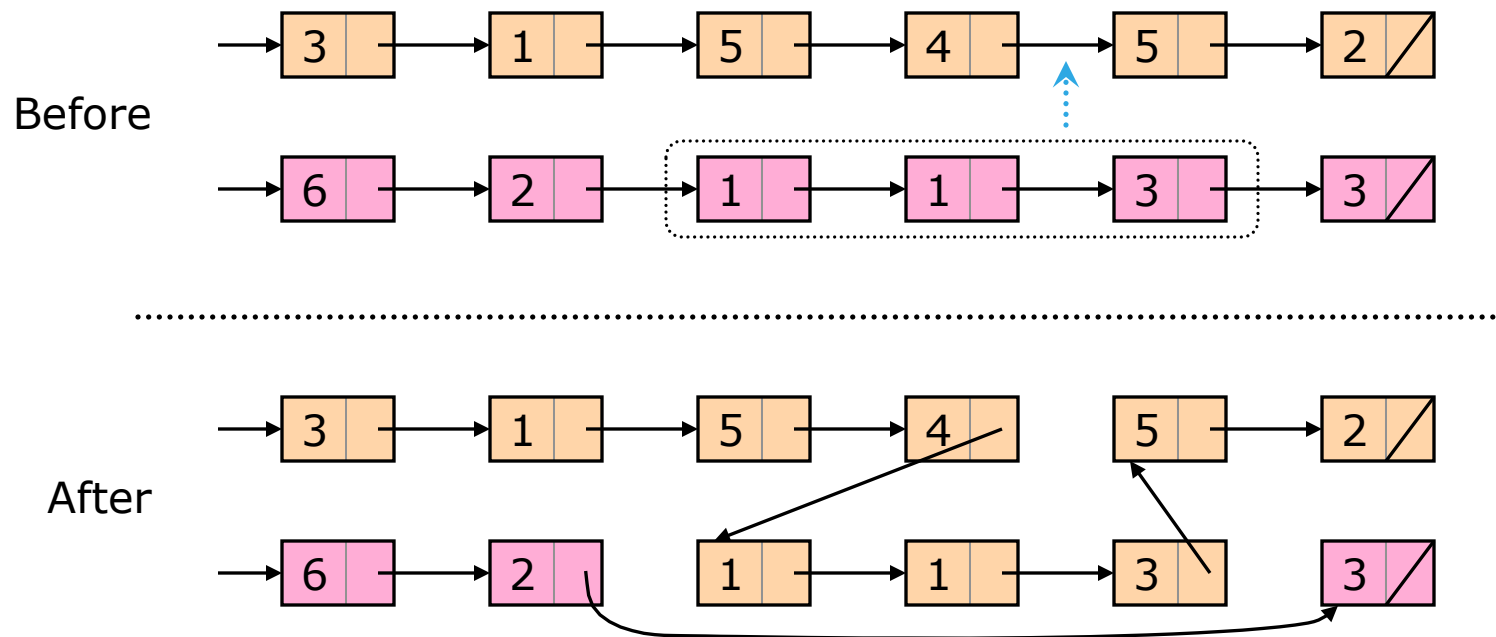
- Removing the item at a given position from an array *is also slow-ish*.
- Removing the item at a given position from a Linked List is very fast.
 - We need an iterator to the *previous* item.
- Example: Remove the item in the bold node.



More on Linked Lists

More Advantages [1/2]

With Linked Lists, we can do a fast **splice**:




Note: If we allow for efficient splicing, then we cannot efficiently keep track of a Linked List's size.

More on Linked Lists

More Advantages [2/2]

With Linked Lists, iterators, pointers, and references to items will always stay valid and never change what they refer to, as long as the Linked List exists — unless we remove or change the item ourselves.



More on Linked Lists

Comparison With Arrays [1/4]

1. What is the order of each of the following when using (a) a smart-array implementation of a Sequence, and (b) a Linked-List implementation? Assume good design and good algorithms.
 - Look-up an item by index.
 - Search in a sorted Sequence.
 - Search in an unsorted Sequence.
 - Sort a Sequence.
 - Insert an item at a given position.
 - Remove an item at a given position.
 - Splice part of one Sequence into another.
 - Insert item at beginning of Sequence.
 - Remove item at beginning of Sequence.
 - Insert item at end of Sequence.
 - Remove item at end of Sequence.
 - Traverse a Sequence (iterate through all items, in order).
 - Make a copy of an entire Sequence.
 - Swap two Sequences.
2. What other issues arise when comparing the two data structures?

More on Linked Lists

Comparison With Arrays [2/4]

	Smart Array	Linked List
Look-up by index	$O(1)$	$O(n)$
Search sorted	$O(\log n)$	$O(n)$
Search unsorted	$O(n)$	$O(n)$
Sort	$O(n \log n)$	$O(n \log n)$
Insert @ given pos	$O(n)$	$O(1)^*$
Remove @ given pos	$O(n)$	$O(1)^*$
Splice	$O(n)$	$O(1)$
Insert @ beginning	$O(n)$	$O(1)$
Remove @ beginning	$O(n)$	$O(1)$
Insert @ end	$O(1)$ or $O(n)^{**}$ amortized const	$O(1)$ or $O(n)^{***}$
Remove @ end	$O(1)$	$O(1)$ or $O(n)^{***}$
Traverse	$O(n)$	$O(n)$
Copy	$O(n)$	$O(n)$
Swap	$O(1)$	$O(1)$

Find faster
with an array

Rearrange faster
with a Linked List

*For Singly Linked Lists, we mean inserting or removing just *after* the given position.

- Doubly Linked Lists can help.

** $O(n)$ if reallocation occurs. Otherwise, $O(1)$. Amortized constant time.

- Pre-allocation can help.

***For $O(1)$, need a pointer to the end of the list. Otherwise, $O(n)$.

- This is tricky.
- And, for remove @ end, it is basically impossible.
- Doubly Linked Lists can help.

More on Linked Lists

Comparison With Arrays [3/4]

Other Issues

- ☹ Linked Lists use **more memory**.
- ☹ When order is the same, Linked Lists are almost always **slower**.
 - Arrays might be 2–10 times faster.
- ☹ Arrays keep consecutive items in **nearby memory locations**.
 - Many algorithms have the property that when they access a data item, the following accesses are likely to be to the same or nearby items.
 - This property of an algorithm is called **locality of reference**.
 - Once a memory location is accessed, the CPU cache can **prefetch** nearby memory locations. With an array, these are likely to hold nearby data items.
 - Thus, because of cache prefetching, an array can have a significant speed advantage over a Linked List, when used with an algorithm that has good locality of reference.
- ☺ With an array, iterators, pointers, and references to items can be **invalidated** by reallocation. Also, insert/remove can change the item they reference.

More on Linked Lists


Comparison With Arrays [4/4]

The Moral of the Story

- Two kinds of design decisions affect the efficiency of code:
 - Choice of algorithm.
 - Choice of data structure.
- The latter often has the greater impact.

Again:


- Use arrays when we want fast look-up/search.
- Use Linked Lists when we want fast insert & delete (by iterator).



More on Linked Lists


Variations — Doubly Linked List [1/3]

The kind of Linked List we have been discussing contains one pointer per node. Thus, it is called a **Singly Linked List**.



In a **Doubly Linked List**, each node has a data item & **two pointers**:

- A pointer to the next node.
- A pointer to the previous node.



Doubly Linked Lists often have an end-of-list pointer.

- This can be efficiently maintained, resulting in constant-time insert and remove at the end.

More on Linked Lists

Variations — Doubly Linked List [2/3]

Doubly Linked Lists have essentially all the advantages of (Singly) Linked Lists, plus some more.

- They allow efficient insert/remove at both ends of the list.
 - We can maintain an end-of-list pointer without trouble.
- They allow efficient insert-before-this-node and remove-this-node.
- They allow efficient backwards iteration.

However

- Doubly Linked Lists are a *little* slower than Singly Linked Lists.
 - Constant-time operations remain $O(1)$, but the constant is a little larger.
- Doubly Linked Lists still cannot do both splice and size efficiently.

The Bottom Line

- Doubly Linked Lists are generally considered to be a good basis for a **general-purpose** generic container type.
 - Singly-Linked Lists are not. Remember all those asterisks?

The C++ STL uses Doubly Linked Lists: the `std::list` class template. (C++11 also has `std::forward_list`, a singly linked list.)

More on Linked Lists

Variations — Doubly Linked List [3/3]

	Smart Array	Doubly Linked List
Look-up by index	$O(1)$	$O(n)$
Search sorted	$O(\log n)$	$O(n)$
Search unsorted	$O(n)$	$O(n)$
Sort	$O(n \log n)$	$O(n \log n)$
Insert @ given pos	$O(n)$	$O(1)$
Remove @ given pos	$O(n)$	$O(1)$
Splice	$O(n)$	$O(1)$
Insert @ beginning	$O(n)$	$O(1)$
Remove @ beginning	$O(n)$	$O(1)$
Insert @ end	$O(1)$ or $O(n)^*$ amortized const	$O(1)$
Remove @ end	$O(1)$	$O(1)$
Traverse	$O(n)$	$O(n)$
Copy	$O(n)$	$O(n)$
Swap	$O(1)$	$O(1)$

With Doubly Linked Lists, we can get rid of most of our asterisks.

* $O(n)$ if reallocation occurs. Otherwise, $O(1)$. Amortized constant time.

- Pre-allocation can help.

Find faster
with an array

Rearrange faster
with a Linked List


More on Linked Lists

Variations — Circular Linked List

Using nodes and pointers, we can arrange data structures in just about any way we want.

An interesting variation on a Linked List is a **Circular Linked List**.

- Here, we arrange our nodes in a circle.



- Ownership issues get trickier.
 - Destroy the head and ... nothing else gets destroyed?
 - One (somewhat icky) solution: use a **weak reference**.

More on Linked Lists Implementation

Two approaches to implementing a Linked List:

- A Linked List package to be used by others.
- An internal-use Linked List: part of some other package, and not exposed to client code.

How would these be different?

- In particular, what classes might we define in each case?
 - In the first case, many classes: container, node, iterator.
 - In the second case, a node class, but possibly nothing else.

TO DO

- Update the internal-use Singly Linked List begun long ago, to include:
 - Exception-safety information.
 - An insert-at-beginning operation.

Sequences in the C++ STL

Generic Sequence Types — Introduction

The C++ STL has four generic Sequence container types.

- Class template `std::vector`.
 - A “smart array”.
 - Much like the Assignment 5 package, but with more member functions.
- Class template `std::basic_string`.
 - Much like `std::vector`, but aimed at character string operations.
 - Mostly we use `std::string`, which is really `std::basic_string<char>`.
 - Also `std::wstring`, which is really `std::basic_string<std::wchar_t>`.
- Class template `std::list`.
 - A Doubly Linked List.
 - Note: The Standard does not specify implementation. It specifies the semantics and order of operations. These were written with a Doubly Linked List in mind, and a D.L.L. is the usual implementation.
- Class template `std::deque`.
 - Deque stands for **D**ouble-**E**nded **Q**UEue.
 - Say “deck”.
 - Like `std::vector`, but a bit slower. Allows fast insert/remove at both beginning and end.

Sequences in the C++ STL


Generic Sequence Types — `std::deque` [1/4]

We are familiar with smart arrays and Linked Lists. How is `std::deque` implemented?


- There are two big ideas behind it.

First Idea

- A **vector** uses an array in which data are stored at the beginning.



- This gives linear-time insert/remove at beginning, constant-time remove at end, and, if we do it right, amortized-constant-time insert at end.
- What if we store data in the middle? When we reallocate-and-copy, we move our data to the middle of the new array.




- Result: Amortized $O(1)$ insert, and $O(1)$ remove, at **both** ends.

Sequences in the C++ STL

Generic Sequence Types — `std::deque` [2/4]

Second Idea

- Doing reallocate-and-copy for a **vector** requires calling either the copy constructor or copy assignment for **every data item**.
 - For large, complex data items, this can be time-consuming.
- Instead, let our array be an **array of pointers to arrays**, so that reallocate-and-copy only needs to move the pointers.
 - This still lets us keep most of the locality-of-reference advantages of an array, when the data items are small.

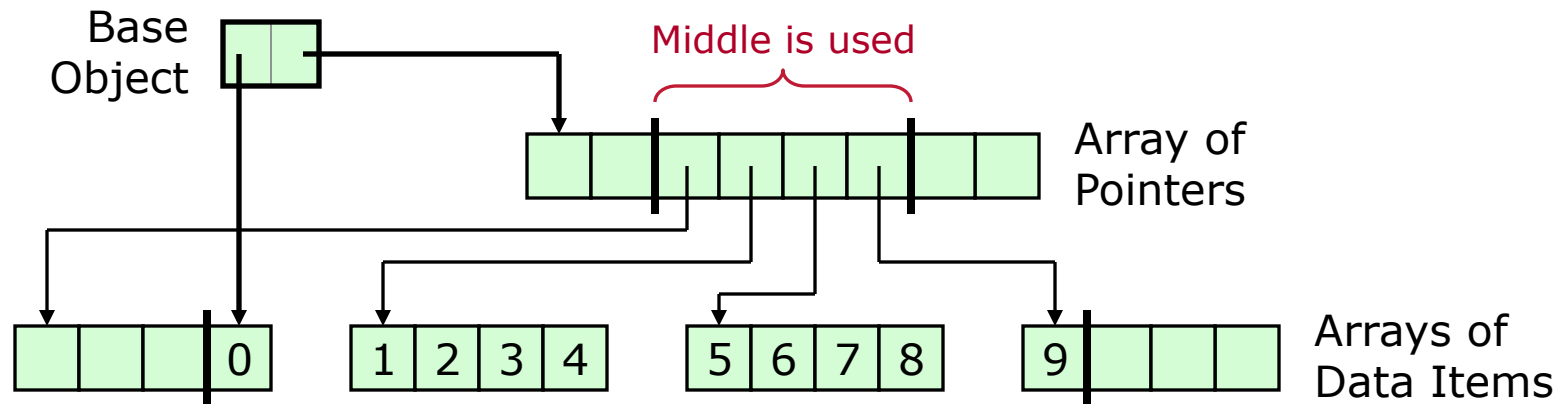


Sequences in the C++ STL

Generic Sequence Types — `std::deque` [3/4]

An implementation of `std::deque` typically uses both of these ideas.

- It probably uses an array of pointers to arrays.
 - This *might* go deeper (array of pointers to arrays of pointers to arrays).
- The arrays may not be filled all the way to the beginning or the end.
- Reallocate-and-copy moves the data to the middle of the new array of pointers.



Thus, `std::deque` is an array-ish container, optimized for:

- Insert/remove at either end.
- Possibly large, difficult-to-copy data items.

The cost is complexity, and a slower [but still $O(1)$] look-up by index.

Sequences in the C++ STL

Generic Sequence Types — `std::deque` [4/4]

Essentially, `std::deque` is an array.

- Iterators are random-access.
- But it has some complexity to it, so it is a slow-ish array.

Like `vector`, `deque` *tends* to keep consecutive items in nearby memory locations.

- So it tends to avoid cache misses, when used with algorithms having good locality of reference.

However:

- Insertions at the beginning do not require items to be moved up.
 - We speed up insert-at-beginning by allocating extra space before existing data.
- Reallocate-and-copy leaves the data items alone.
 - We also speeds up insertion by trading value-type operations for pointer operations.
 - Pointer operations can be much faster than value-type operations. A `std::deque` can do reallocate-and-copy using a raw memory copy, with no value-type copy ctor calls.

The Bottom Line

- A `std::deque` is generally a good choice when you need fast insert/remove at both ends of a Sequence.
- Especially if you also want fast-ish look-up.
- Some people also recommend `std::deque` whenever you will be doing a lot of resizing, but do not need fast insert/remove in the middle.

Sequences in the C++ STL

Generic Sequence Types — Efficiency [1/2]

We measure efficiency by counting steps. How do we count steps for a generic container type?

- We count both built-in operations and value-type operations.
- However, we typically expect that the most time-consuming operations are those on the value type.

The C++ Standard, on the other hand, counts **only** value-type operations.

- For example, “constant time” in the Standard means that at most a constant number of value-type operations are performed.

Sequences in the C++ STL

Generic Sequence Types — Efficiency [2/2]

	vector, basic_string	deque	list
Look-up by index	Constant	Constant	Linear
Search sorted	Logarithmic	Logarithmic	Linear
Insert @ given pos	Linear	Linear	Constant
Remove @ given pos	Linear	Linear	Constant
Insert @ beginning	Linear	Linear/ Amortized Constant*	Constant
Remove @ beginning	Linear	Constant	Constant
Insert @ end	Linear/ Amortized Constant**	Linear/ Amortized Constant*	Constant
Remove @ end	Constant	Constant	Constant

*Only a constant number of value-type operations are required.

- The C++ standard counts only value-type operations. Thus, it says that insert at the beginning or end of a `std::deque` is constant time.

**Constant time if sufficient memory has already been allocated.

All have $O(n)$ traverse, copy, and search-unsorted, $O(1)$ swap, and $O(n \log n)$ sort.

Sequences in the C++ STL

Generic Sequence Types — Common Features

All STL Sequence containers have:

- **iterator, const_iterator**
 - Iterator types. The latter acts like a pointer-to-const.
 - **vector, basic_string, deque** have random-access iterators.
 - **list** has bidirectional iterators.
- **iterator begin(), iterator end()**
- **iterator insert(iterator, item)**
 - Insert before. Returns position of new item.
- **iterator erase(iterator)**
 - Remove this item. Returns position of following item.
- **push_back(item)**
 - Insert at the end.
- **clear()**
 - Remove all items.
- **resize(newSize)**
 - Change the size of the container.
 - Not the same as **vector::reserve**, which sets capacity.

In Addition

vector, deque, list have:

- **pop_back()**
 - Remove at the end.
- **reference front(), reference back()**
 - Return reference to first, last item.

deque, list have:

- **push_front(item), pop_front()**
 - Insert & remove at the beginning.

vector, basic_string, deque have:

- **reference operator[] (index)**
 - Look-up by index.

vector has:

- **reserve(newCapacity)**
 - Sets capacity to at least the given value.

And there are other members ...